


# Customization basics: tensors and operations

Run in

[Google](https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/customization-basics/tensors-and-operations.ipynb) (<https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/customization-basics/tensors-and-operations.ipynb>)  
[Colab](#)



This is an introductory TensorFlow tutorial that shows how to:

- Import the required package
- Create and use tensors
- Use GPU acceleration
- Demonstrate `tf.data.Dataset` ([/api\\_docs/python/tf/data/Dataset](/api_docs/python/tf/data/Dataset))

## Import TensorFlow

To get started, import the `tensorflow` module. As of TensorFlow 2, eager execution is turned on by default. This enables a more interactive frontend to TensorFlow, the details of which we will discuss much later.

```
import tensorflow as tf
```

## Tensors

A Tensor is a multi-dimensional array. Similar to NumPy `ndarray` objects, `tf.Tensor` ([/api\\_docs/python/tf/Tensor](/api_docs/python/tf/Tensor)) objects have a data type and a shape. Additionally, `tf.Tensor` ([/api\\_docs/python/tf/Tensor](/api_docs/python/tf/Tensor))s can reside in accelerator memory (like a GPU). TensorFlow offers a rich library of operations (`tf.add` ([https://www.tensorflow.org/api\\_docs/python/tf/add](https://www.tensorflow.org/api_docs/python/tf/add)), `tf.matmul` ([https://www.tensorflow.org/api\\_docs/python/tf/matmul](https://www.tensorflow.org/api_docs/python/tf/matmul)), `tf.linalg.inv` ([https://www.tensorflow.org/api\\_docs/python/tf/linalg/inv](https://www.tensorflow.org/api_docs/python/tf/linalg/inv)) etc.) [More details](#) and produce

This site uses cookies from Google to deliver its services and to analyze traffic.

[More details](#) [OK](#)

**tf.Tensor** (/api\_docs/python/tf/Tensor)s. These operations automatically convert native Python types, for example:

```
:(tf.add(1, 2))
:(tf.add([1, 2], [3, 4]))
:(tf.square(5))
:(tf.reduce_sum([1, 2, 3]))
```

operator overloading is also supported

```
:(tf.square(2) + tf.square(3))
```

```
tf.Tensor(3, shape=(), dtype=int32)
tf.Tensor([4 6], shape=(2,), dtype=int32)
tf.Tensor(25, shape=(), dtype=int32)
tf.Tensor(6, shape=(), dtype=int32)
tf.Tensor(13, shape=(), dtype=int32)
```

Each **tf.Tensor** (/api\_docs/python/tf/Tensor) has a shape and a datatype:

```
tf.matmul([[1]], [[2, 3]])
x
x.shape
x.dtype
```

```
tf.Tensor([[2 3]], shape=(1, 2), dtype=int32)
)
In: 'int32'>
```

The most obvious differences between NumPy arrays and **tf.Tensor** (/api\_docs/python/tf/Tensor)s are:

1. Tensors can be backed by accelerator memory (like GPU, TPU).

2. Tensors are immutable.

[More details](#) [OK](#)

# NumPy Compatibility

Converting between a TensorFlow `tf.Tensor` (/api\_docs/python/tf/Tensor)s and a NumPy `ndarray` is easy:

- TensorFlow operations automatically convert NumPy `ndarrays` to Tensors.
- NumPy operations automatically convert Tensors to NumPy `ndarrays`.

Tensors are explicitly converted to NumPy `ndarrays` using their `.numpy()` method. These conversions are typically cheap since the array and `tf.Tensor` (/api\_docs/python/tf/Tensor) share the underlying memory representation, if possible. However, sharing the underlying representation isn't always possible since the `tf.Tensor` (/api\_docs/python/tf/Tensor) may be hosted in GPU memory while NumPy arrays are always backed by host memory, and the conversion involves a copy from GPU to host memory.

```
import numpy as np

x = np.ones([3, 3])

# TensorFlow operations convert numpy arrays to Tensors automatically
y = tf.multiply(ndarray, 42)
print(y)

# And NumPy operations convert Tensors to numpy arrays automatically
print(np.add(tensor, 1))

# The .numpy() method explicitly converts a Tensor to a numpy array
print(tensor.numpy())
```

```
TensorFlow operations convert numpy arrays to Tensors automatically
tensor(
  42. 42.]
  42. 42.]
  42. 42.]], shape=(3, 3), dtype=float64)
NumPy operations convert Tensors to numpy arrays automatically
43. 43.]
43. 43.]
43. 43.]]
```

`.numpy()` method explicitly converts a Tensor to a numpy array

```
42. 42.]
42. 42.]
```

[More details](#) [OK](#)

## GPU acceleration

Many TensorFlow operations are accelerated using the GPU for computation. Without any annotations, TensorFlow automatically decides whether to use the GPU or CPU for an operation—copying the tensor between CPU and GPU memory, if necessary. Tensors produced by an operation are typically backed by the memory of the device on which the operation executed, for example:

```
tf.random.uniform([3, 3])

("Is there a GPU available: "),
(tf.config.experimental.list_physical_devices("GPU"))

("Is the Tensor on GPU #0: "),
(x.device.endswith('GPU:0'))

here a GPU available:
icalDevice(name='/physical_device:GPU:0', device_type='GPU')]
ie Tensor on GPU #0:
```

## Device Names

The **Tensor.device** ([/api\\_docs/python/tf/Tensor#device](/api_docs/python/tf/Tensor#device)) property provides a fully qualified string name of the device hosting the contents of the tensor. This name encodes many details, such as an identifier of the network address of the host on which this program is executing and the device within that host. This is required for distributed execution of a TensorFlow program. The string ends with **GPU:<N>** if the tensor is placed on the **N**-th GPU on the host.

## Explicit Device Placement

This site uses cookies from Google to deliver its services and to analyze traffic.  
In TensorFlow, *placement* refers to how individual operations are assigned (placed on) a device for execution. As mentioned, when there is no explicit **More details** provided, TensorFlow automatically decides which device to execute an operation and copies tensors

to that device, if needed. However, TensorFlow operations can be explicitly placed on specific devices using the `tf.device` ([/api\\_docs/python/tf/device](/api_docs/python/tf/device)) context manager, for example:

```
t time

time_matmul(x):
    start = time.time()
    for loop in range(10):
        tf.matmul(x, x)

    result = time.time()-start

    print("10 loops: {:.2f}ms".format(1000*result))

# Execution on CPU
print("On CPU:")
with tf.device("CPU:0"):
    x = tf.random.uniform([1000, 1000])
    assert x.device.endswith("CPU:0")
    time_matmul(x)

# Execution on GPU #0 if available
tf.config.experimental.list_physical_devices("GPU"):
print("On GPU:")
with tf.device("GPU:0"): # Or GPU:1 for the 2nd GPU, GPU:2 for the 3rd etc.
    x = tf.random.uniform([1000, 1000])
    assert x.device.endswith("GPU:0")
    time_matmul(x)

CPU:
10 loops: 96.68ms
GPU:
10 loops: 271.75ms
```

## Datasets

This site uses cookies from Google to deliver its services and to analyze traffic.

This section uses the `tf.data.Dataset` API (<https://www.tensorflow.org/guide/datasets>) to build a pipeline for feeding data to your model. The `tf.data.Dataset`

[More details](#) 

([/api\\_docs/python/tf/data/Dataset](https://www.tensorflow.org/api_docs/python/tf/data/Dataset)) API is used to build performant, complex input pipelines from simple, re-usable pieces that will feed your model's training or evaluation loops.

## Create a source Dataset

Create a *source* dataset using one of the factory functions like [Dataset.from\\_tensors](https://www.tensorflow.org/api_docs/python/tf/data/Dataset#from_tensors) ([https://www.tensorflow.org/api\\_docs/python/tf/data/Dataset#from\\_tensors](https://www.tensorflow.org/api_docs/python/tf/data/Dataset#from_tensors)),

[Dataset.from\\_tensor\\_slices](https://www.tensorflow.org/api_docs/python/tf/data/Dataset#from_tensor_slices)

([https://www.tensorflow.org/api\\_docs/python/tf/data/Dataset#from\\_tensor\\_slices](https://www.tensorflow.org/api_docs/python/tf/data/Dataset#from_tensor_slices)), or using objects that read from files like [TextLineDataset](https://www.tensorflow.org/api_docs/python/tf/data/TextLineDataset)

([https://www.tensorflow.org/api\\_docs/python/tf/data/TextLineDataset](https://www.tensorflow.org/api_docs/python/tf/data/TextLineDataset)) or [TFRecordDataset](https://www.tensorflow.org/api_docs/python/tf/data/TFRecordDataset) ([https://www.tensorflow.org/api\\_docs/python/tf/data/TFRecordDataset](https://www.tensorflow.org/api_docs/python/tf/data/TFRecordDataset)). See the [TensorFlow Dataset guide](https://www.tensorflow.org/guide/datasets#reading_input_data) ([https://www.tensorflow.org/guide/datasets#reading\\_input\\_data](https://www.tensorflow.org/guide/datasets#reading_input_data)) for more information.

```
ensors = tf.data.Dataset.from_tensor_slices([1, 2, 3, 4, 5, 6])
```

```
ate a CSV file
```

```
t tempfile
```

```
lename = tempfile.mkstemp()
```

```
open(filename, 'w') as f:
```

```
rite("""Line 1
```

```
2
```

```
3
```

```
)
```

```
le = tf.data.TextLineDataset(filename)
```

## Apply transformations

Use the transformations functions like [map](https://www.tensorflow.org/api_docs/python/tf/data/Dataset#map)

([https://www.tensorflow.org/api\\_docs/python/tf/data/Dataset#map](https://www.tensorflow.org/api_docs/python/tf/data/Dataset#map)), [batch](https://www.tensorflow.org/api_docs/python/tf/data/Dataset#batch)

([https://www.tensorflow.org/api\\_docs/python/tf/data/Dataset#batch](https://www.tensorflow.org/api_docs/python/tf/data/Dataset#batch)), and [shuffle](https://www.tensorflow.org/api_docs/python/tf/data/Dataset#shuffle)

([https://www.tensorflow.org/api\\_docs/python/tf/data/Dataset#shuffle](https://www.tensorflow.org/api_docs/python/tf/data/Dataset#shuffle)) to apply transformations to dataset records.

This site uses cookies from Google to deliver its services and to analyze traffic.

```
ensors = ds_tensors.map(tf.square).shuffle(2).batch(2)
```

[More details](#)

[OK](#)

```
le = ds_file.batch(2)
```

## Iterate

**`tf.data.Dataset`** ([/api\\_docs/python/tf/data/Dataset](https://api_docs/python/tf/data/Dataset)) objects support iteration to loop over records:

```
for (Elements of ds_tensors:)  
    for in ds_tensors:  
        print(x)
```

```
for (\nElements in ds_file:)  
    for in ds_file:  
        print(x)
```

```
for tensors of ds_tensors:  
    print(tensor([1 9], shape=(2,), dtype=int32))  
    print(tensor([16 25], shape=(2,), dtype=int32))  
    print(tensor([ 4 36], shape=(2,), dtype=int32))
```

```
for tensors in ds_file:  
    print(tensor([b'Line 1' b'Line 2'], shape=(2,), dtype=string))  
    print(tensor([b'Line 3' b'   '], shape=(2,), dtype=string))
```

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (<https://creativecommons.org/licenses/by/4.0/>), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (<https://www.apache.org/licenses/LICENSE-2.0>). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies) (<https://developers.google.com/site-policies>). Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2020-06-12.

This site uses cookies from Google to deliver its services and to analyze traffic.

[More details](#) [OK](#)