

Distributed training with Keras

Run in

[Google](https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/dist) (<https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/dist>)

[Colab](#)



Overview

The `tf.distribute.Strategy` (/api_docs/python/tf/distribute/Strategy) API provides an abstraction for distributing your training across multiple processing units. The goal is to allow users to enable distributed training using existing models and training code, with minimal changes.

This tutorial uses the `tf.distribute.MirroredStrategy` (/api_docs/python/tf/distribute/MirroredStrategy), which does in-graph replication with synchronous training on many GPUs on one machine. Essentially, it copies all of the model's variables to each processor. Then, it uses `all-reduce` (<http://mpitutorial.com/tutorials/mpi-reduce-and-allreduce/>) to combine the gradients from all processors and applies the combined value to all copies of the model.

`MirroredStrategy` is one of several distribution strategy available in TensorFlow core. You can read about more strategies at [distribution strategy guide](https://www.tensorflow.org/guide/distributed_training) (https://www.tensorflow.org/guide/distributed_training).

Keras API

This example uses the `tf.keras` (/api_docs/python/tf/keras) API to build the model and training loop. For custom training loops, see the [tf.distribute.Strategy with training loops](/tutorials/distribute/training_loops) (/tutorials/distribute/training_loops) tutorial.

Import dependencies

This site uses cookies from Google to deliver its services and to analyze traffic.

ort TensorFlow and TensorFlow Datasets

[More details](#)

OK

```

import tensorflow_datasets as tfds
import tensorflow as tf
tf.disable_progress_bar()

import os

```

```

print(tf.__version__)

```

```

)

```

Download the dataset

Download the MNIST dataset and load it from [TensorFlow Datasets](https://www.tensorflow.org/datasets) (<https://www.tensorflow.org/datasets>). This returns a dataset in [tf.data](https://api_docs/python/tf/data) ([/api_docs/python/tf/data](https://api_docs/python/tf/data)) format.

Setting `with_info` to `True` includes the metadata for the entire dataset, which is being saved here to `info`. Among other things, this metadata object includes the number of train and test examples.

```

datasets, info = tfds.load(name='mnist', with_info=True, as_supervised=True)

mnist_train, mnist_test = datasets['train'], datasets['test']

```

Define distribution strategy

Create a `MirroredStrategy` object. This will handle distribution, and provides a context manager ([tf.distribute.MirroredStrategy.scope](https://api_docs/python/tf/distribute/MirroredStrategy#scope) ([/api_docs/python/tf/distribute/MirroredStrategy#scope](https://api_docs/python/tf/distribute/MirroredStrategy#scope))) to build your model inside.

This site uses cookies from Google to deliver its services and to analyze traffic.

```

strategy = tf.distribute.MirroredStrategy()

```

[More details](#)

OK

```
tensorflow:Using MirroredStrategy with devices ('/job:localhost/replica:0/task:0/device:GPU:0')
tensorflow:Using MirroredStrategy with devices ('/job:localhost/replica:0/task:0/device:GPU:0')
```

```
('Number of devices: {}'.format(strategy.num_replicas_in_sync))
```

```
Number of devices: 1
```

Setup input pipeline

When training a model with multiple GPUs, you can use the extra computing power effectively by increasing the batch size. In general, use the largest batch size that fits the GPU memory, and tune the learning rate accordingly.

You can also do `info.split_info.total_num_examples` to get the total number of examples in the dataset.

```
train_examples = info.split_info['train'].num_examples
test_examples = info.split_info['test'].num_examples
```

```
BATCH_SIZE = 10000
```

```
BATCH_SIZE_PER_REPLICA = 64
```

```
BATCH_SIZE = BATCH_SIZE_PER_REPLICA * strategy.num_replicas_in_sync
```

Pixel values, which are 0-255, have to be normalized to the 0-1 range (https://en.wikipedia.org/wiki/Feature_scaling). Define this scale in a function.

```
def scale_image(image, label):
    image = tf.cast(image, tf.float32)
    image /= 255
```

[More details](#) [OK](#)

```
return image, label
```

Apply this function to the training and test data, shuffle the training data, and batch it for training (https://www.tensorflow.org/api_docs/python/tf/data/Dataset#batch). Notice we are also keeping an in-memory cache of the training data to improve performance.

```
train_dataset = mnist_train.map(scale).cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)
test_dataset = mnist_test.map(scale).batch(BATCH_SIZE)
```

Create the model

Create and compile the Keras model in the context of `strategy.scope`.

```
strategy.scope():
    model = tf.keras.Sequential([
        tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=(28, 28, 1)),
        tf.keras.layers.MaxPooling2D(),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(64, activation='relu'),
        tf.keras.layers.Dense(10)

    ])

model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              optimizer=tf.keras.optimizers.Adam(),
              metrics=['accuracy'])
```

Define the callbacks

The callbacks used here are:

- *TensorBoard*: This callback writes a log for TensorBoard which allows you to visualize the graphs.
- *Model Checkpoint*: This callback saves the model after every epoch.
- *Learning Rate Scheduler*: Using this callback, you can schedule the learning rate to change after every epoch/batch.

[More details](#) [OK](#)

For illustrative purposes, add a print callback to display the *learning rate* in the notebook.

```

line the checkpoint directory to store the checkpoints

checkpoint_dir = './training_checkpoints'
# Name of the checkpoint files
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt_{epoch}")

# Callback for decaying the learning rate.
# You can define any decay function you need.
def decay(epoch):
    if epoch < 3:
        return 1e-3
    elif epoch >= 3 and epoch < 7:
        return 1e-4
    else:
        return 1e-5

# Callback for printing the LR at the end of each epoch.
class PrintLR(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs=None):
        print('\nLearning rate for epoch {} is {}'.format(epoch + 1,
                                                            model.optimizer.lr.numpy()))

callbacks = [
    tf.keras.callbacks.TensorBoard(log_dir='./logs'),
    tf.keras.callbacks.ModelCheckpoint(filepath=checkpoint_prefix,
                                       save_weights_only=True),
    tf.keras.callbacks.LearningRateScheduler(decay),
    PrintLR()
]

```

Train and evaluate

Now, train the model in the usual way, calling `fit` on the model and passing in the dataset created at the beginning of the tutorial. This step is the same whether you are distributing the training or not.

[More details](#) [OK](#)

```
..fit(train_dataset, epochs=12, callbacks=callbacks)
```

```
| 1/12
```

```
tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadc  
tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadc  
tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadc  
tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadc  
tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadc  
tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadc
```

As you can see below, the checkpoints are getting saved.

```
ck the checkpoint directory  
checkpoint_dir}
```

```
point          ckpt_4.data-00000-of-00002  
.1.data-00000-of-00002  ckpt_4.data-00001-of-00002  
.1.data-00001-of-00002  ckpt_4.index  
.1.index             ckpt_5.data-00000-of-00002  
.10.data-00000-of-00002 ckpt_5.data-00001-of-00002  
.10.data-00001-of-00002 ckpt_5.index  
.10.index            ckpt_6.data-00000-of-00002  
.11.data-00000-of-00002 ckpt_6.data-00001-of-00002  
.11.data-00001-of-00002 ckpt_6.index  
.11.index            ckpt_7.data-00000-of-00002  
.12.data-00000-of-00002 ckpt_7.data-00001-of-00002  
.12.data-00001-of-00002 ckpt_7.index  
.12.index            ckpt_8.data-00000-of-00002  
.2.data-00000-of-00002  ckpt_8.data-00001-of-00002
```

To see how the model perform, load the latest checkpoint and call **evaluate** on the test data.

This site uses cookies from Google to deliver its services and to analyze traffic.

[More details](#) [OK](#)

Call **evaluate** as before using appropriate datasets.

```
..load_weights(tf.train.latest_checkpoint(checkpoint_dir))

loss, eval_acc = model.evaluate(eval_dataset)

:('Eval loss: {}, Eval Accuracy: {}'.format(eval_loss, eval_acc))
```

```
tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadc
tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadc
tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadc
tensorflow:Reduce to /job:localhost/replica:0/task:0/device:CPU:0 then broadc

57 [=====] - 1s 6ms/step - accuracy: 0.9859 - loss:
loss: 0.040106043219566345, Eval Accuracy: 0.9858999848365784
```

To see the output, you can download and view the TensorBoard logs at the terminal.

```
tensorboard --logdir=path/to/log-directory
```

```
-sh ./logs
```

```
. 4.0K
train
```

Export to SavedModel

This site uses cookies from Google to deliver its services and to analyze traffic.
Export the graph and the variables to the platform-agnostic SavedModel format. After your model is saved, you can load it with or without the scope. [More details](#) [OK](#)

```
= 'saved_model/'
```

```
..save(path, save_format='tf')
```

```
INFO:tensorflow:From /tmpfs/src/tf_docs_env/lib/python3.6/site-packages/tensorflow/core/framework/ops.py:510:
tf.nn.conv2d (from tensorflow.nn.conv2d) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.nn.conv2d instead.
INFO:tensorflow:From /tmpfs/src/tf_docs_env/lib/python3.6/site-packages/tensorflow/core/framework/ops.py:510:
tf.nn.conv2d (from tensorflow.nn.conv2d) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.nn.conv2d instead.
INFO:tensorflow:From /tmpfs/src/tf_docs_env/lib/python3.6/site-packages/tensorflow/core/framework/ops.py:510:
tf.nn.conv2d (from tensorflow.nn.conv2d) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.nn.conv2d instead.
INFO:tensorflow:From /tmpfs/src/tf_docs_env/lib/python3.6/site-packages/tensorflow/core/framework/ops.py:510:
tf.nn.conv2d (from tensorflow.nn.conv2d) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.nn.conv2d instead.
INFO:tensorflow:Assets written to: saved_model/assets
```

```
INFO:tensorflow:From /tmpfs/src/tf_docs_env/lib/python3.6/site-packages/tensorflow/core/framework/ops.py:510:
tf.nn.conv2d (from tensorflow.nn.conv2d) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.nn.conv2d instead.
INFO:tensorflow:From /tmpfs/src/tf_docs_env/lib/python3.6/site-packages/tensorflow/core/framework/ops.py:510:
tf.nn.conv2d (from tensorflow.nn.conv2d) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.nn.conv2d instead.
INFO:tensorflow:Assets written to: saved_model/assets
```

```
tensorflow:Assets written to: saved_model/assets
```

```
tensorflow:Assets written to: saved_model/assets
```

Load the model without `strategy.scope`.

```
unreplicated_model = tf.keras.models.load_model(path)

unreplicated_model.compile(
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    optimizer=tf.keras.optimizers.Adam(),
    metrics=['accuracy'])

eval_loss, eval_acc = unreplicated_model.evaluate(eval_dataset)

print('Eval loss: {}, Eval Accuracy: {}'.format(eval_loss, eval_acc))
```

```
57 [=====] - 1s 5ms/step - loss: 0.0401 - accuracy: 0.9858999848365784
```

This site uses cookies from Google to deliver its services and to analyze traffic.

[More details](#) [OK](#)

Load the model with `strategy.scope`.


```

strategy.scope():
replicated_model = tf.keras.models.load_model(path)
replicated_model.compile(loss=tf.keras.losses.SparseCategoricalCrossentropy(from
optimizer=tf.keras.optimizers.Adam(),
metrics=[ 'accuracy' ]))

val_loss, eval_acc = replicated_model.evaluate(eval_dataset)
print ('Eval loss: {}, Eval Accuracy: {}'.format(eval_loss, eval_acc))

```

```

57 [=====] - 1s 6ms/step - accuracy: 0.9859 - loss:
loss: 0.040106043219566345, Eval Accuracy: 0.9858999848365784

```

Examples and Tutorials

Here are some examples for using distribution strategy with keras fit/compile:

1. Transformer

(https://github.com/tensorflow/models/blob/master/official/nlp/transformer/transformer_main.py)

example trained using **tf.distribute.MirroredStrategy**

([/api_docs/python/tf/distribute/MirroredStrategy](#))

2. NCF

(https://github.com/tensorflow/models/blob/master/official/recommendation/ncf_keras_main.py)

example trained using **tf.distribute.MirroredStrategy**

([/api_docs/python/tf/distribute/MirroredStrategy](#)).

More examples listed in the [Distribution strategy guide](#)

(https://www.tensorflow.org/guide/distributed_training#examples_and_tutorials)

Next steps

This site uses cookies from Google to deliver its services and to analyze traffic.

- Read the [distribution strategy guide](https://www.tensorflow.org/guide/distributed_training) (https://www.tensorflow.org/guide/distributed_training).
- Read the [Distributed Training with Custom Training Loops](#) **More details** **OK**
([/tutorials/distribute/training_loops](#)) tutorial.

- Visit the [Performance section](https://www.tensorflow.org/guide/function) (<https://www.tensorflow.org/guide/function>) in the guide to learn more about other strategies and [tools](https://www.tensorflow.org/guide/profiler) (<https://www.tensorflow.org/guide/profiler>) you can use to optimize the performance of your TensorFlow models.

`tf.distribute.Strategy` (/api_docs/python/tf/distribute/Strategy) is actively under development and will bring more examples and tutorials in the near future. Please give it a try. We welcome your feedback via [issues](https://github.com/tensorflow/tensorflow/issues/new) (<https://github.com/tensorflow/tensorflow/issues/new>).

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (<https://creativecommons.org/licenses/by/4.0/>), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (<https://www.apache.org/licenses/LICENSE-2.0>). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies) (<https://developers.google.com/site-policies>). Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2020-06-12.

This site uses cookies from Google to deliver its services and to analyze traffic.

[More details](#) [OK](#)