

TensorFlow 2 quickstart for experts

Run in

[Google](https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/quickstart/quickstart.ipynb) (https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/quickstart/quickstart.ipynb)

[Colab](#)

This is a [Google Colaboratory](https://colab.research.google.com/notebooks/welcome.ipynb) notebook file. Python programs are run directly in the browser—a great way to learn and use TensorFlow. To follow this tutorial, run the notebook in Google Colab by clicking the button at the top of this page.

1. In Colab, connect to a Python runtime: At the top-right of the menu bar, select *CONNECT*.
2. Run all the notebook code cells: Select *Runtime > Run all*.

Download and install TensorFlow 2. Import TensorFlow into your program:

Upgrade **pip** to install the TensorFlow 2 package. See the [install guide](https://www.tensorflow.org/install) (https://www.tensorflow.org/install).

Import TensorFlow into your program:

```
import tensorflow as tf

from tensorflow.keras.layers import Dense, Flatten, Conv2D
from tensorflow.keras import Model
```

Load and prepare the [MNIST dataset](http://yann.lecun.com/exdb/mnist/) (http://yann.lecun.com/exdb/mnist/).

```
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

# Add a channels dimension
x_train = x_train[..., tf.newaxis]
x_test = x_test[..., tf.newaxis]
```

[More details](#) [OK](#)

Use `tf.data` (/api_docs/python/tf/data) to batch and shuffle the dataset:

```
train_ds = tf.data.Dataset.from_tensor_slices(
    (x_train, y_train)).shuffle(10000).batch(32)

test_ds = tf.data.Dataset.from_tensor_slices((x_test, y_test)).batch(32)
```

Build the `tf.keras` (/api_docs/python/tf/keras) model using the Keras [model subclassing API](https://www.tensorflow.org/guide/keras#model_subclassing) (https://www.tensorflow.org/guide/keras#model_subclassing):

```
class MyModel(Model):
    def __init__(self):
        super(MyModel, self).__init__()
        self.conv1 = Conv2D(32, 3, activation='relu')
        self.flatten = Flatten()
        self.d1 = Dense(128, activation='relu')
        self.d2 = Dense(10)

    def call(self, x):
        x = self.conv1(x)
        x = self.flatten(x)
        x = self.d1(x)
        return self.d2(x)

# Create an instance of the model
model = MyModel()
```

Choose an optimizer and loss function for training:

```
loss_object = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)

optimizer = tf.keras.optimizers.Adam()
```

Select metrics to measure the loss and the accuracy of the model. These metrics accumulate the values over epochs and then print the overall result.

This site uses cookies from Google to deliver its services and to analyze traffic.

```
train_loss = tf.keras.metrics.Mean(name='train_loss') More details OK
train_accuracy = tf.keras.metrics.SparseCategoricalAccuracy(name='train_accuracy')
```

```
loss = tf.keras.metrics.Mean(name='test_loss')
accuracy = tf.keras.metrics.SparseCategoricalAccuracy(name='test_accuracy')
```

Use [tf.GradientTape](/api_docs/python/tf/GradientTape) (/api_docs/python/tf/GradientTape) to train the model:

```
def train_step(images, labels):
    with tf.GradientTape() as tape:
        # training=True is only needed if there are layers with different
        # behavior during training versus inference (e.g. Dropout).
        predictions = model(images, training=True)
        loss = loss_object(labels, predictions)
        gradients = tape.gradient(loss, model.trainable_variables)
        optimizer.apply_gradients(zip(gradients, model.trainable_variables))

    train_loss(loss)
    train_accuracy(labels, predictions)
```

Test the model:

```
def test_step(images, labels):
    # training=False is only needed if there are layers with different
    # behavior during training versus inference (e.g. Dropout).
    predictions = model(images, training=False)
    loss = loss_object(labels, predictions)

    test_loss(test_loss)
    test_accuracy(labels, predictions)
```

EPOCHS = 5

```
for epoch in range(EPOCHS):
    # reset the metrics at the start of the next epoch
    train_loss.reset_states()
    train_accuracy.reset_states()
    test_loss.reset_states()
    test_accuracy.reset_states()
```

```
for images, labels in train_ds:
    train_step(images, labels)
```

[More details](#) [OK](#)

```

    test_images, test_labels in test_ds:
        test_step(test_images, test_labels)

template = 'Epoch {}, Loss: {}, Accuracy: {}, Test Loss: {}, Test Accuracy: {}'
print(template.format(epoch + 1,
                      train_loss.result(),
                      train_accuracy.result() * 100,
                      test_loss.result(),
                      test_accuracy.result() * 100))

```

WARNING:tensorflow:Layer my_model is casting an input tensor from dtype float64 to dtype float32 to be compatible with the layer. If you intended to run this layer in float32, you can safely ignore this warning.

To change all layers to have dtype float64 by default, call `tf.keras.backend.set_floatx('float64')`.

```

1, Loss: 0.13576626777648926, Accuracy: 95.89666748046875, Test Loss: 0.0658
2, Loss: 0.04346510395407677, Accuracy: 98.66667175292969, Test Loss: 0.0584
3, Loss: 0.024708667770028114, Accuracy: 99.1883316040039, Test Loss: 0.0513
4, Loss: 0.014630611054599285, Accuracy: 99.49500274658203, Test Loss: 0.056
5, Loss: 0.010477297939360142, Accuracy: 99.61333465576172, Test Loss: 0.064

```

The image classifier is now trained to ~98% accuracy on this dataset. To learn more, read the [TensorFlow tutorials](https://www.tensorflow.org/tutorials) (<https://www.tensorflow.org/tutorials>).

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (<https://creativecommons.org/licenses/by/4.0/>), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (<https://www.apache.org/licenses/LICENSE-2.0>). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies) (<https://developers.google.com/site-policies>). Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2020-06-12.

This site uses cookies from Google to deliver its services and to analyze traffic.

[More details](#) [OK](#)